

Framework for GNNs on TUDataset

Erik Deinzer and Felix Glatzel

July 15, 2025

Project Report for

Prof. Rizzi

Prof. De Santis

Sapienza University of Rome

July 15, 2025

Contents

1	Introduction	3
2	Methods	3
2.1	Datasets	3
2.2	Summary Table	7
2.3	Graph Neural Network Architectures	7
2.4	Graph Convolutional Network (GCN)	7
2.5	Graph Isomorphism Network (GIN)	8
2.6	Graph Attention Network (GAT)	9
2.7	GraphSAGE	10
3	Training Setup	11
3.1	Data Preprocessing	11
3.2	Model Architecture	11
3.3	Optimization and Training Strategy	12
3.4	Cross-Validation and Early Stopping	12
3.5	Evaluation Metric	12
4	Results and Discussion	12
4.1	Performance of GNNs Across Datasets	12
4.2	Model Comparison on Individual Datasets	15
4.3	General Observations	16

1 Introduction

Graph-structured data is increasingly prevalent across a wide range of domains, including chemistry, biology, social networks, and recommendation systems. Traditional machine learning techniques often struggle to handle such data due to its non-Euclidean nature. Graph Neural Networks (GNNs) have emerged as a powerful class of models designed to learn representations from graph data by leveraging the underlying structure of nodes and edges.

This project explores the application of GNNs to graph classification tasks using real-world datasets from the TUDataset collection. Specifically, we design and implement a modular and extensible pipeline capable of loading, preprocessing, training, and evaluating models on multiple datasets, including MUTAG, PROTEINS, and ENZYMES. These datasets are commonly used benchmarks in the GNN research community and represent various challenges in terms of graph size, sparsity, heterogeneity, and domain-specific semantics.

Our objectives are threefold: (1) to deepen our understanding of how GNNs can be applied to graph classification problems; (2) to gain practical experience with PyTorch Geometric (PyG) and the TUDataset API; and (3) to analyze and compare the performance of different GNN architectures across datasets with diverse structural characteristics.

In addition to implementing a working framework, we provide a detailed analysis of the structure and nature of each dataset, including an examination of what the nodes and edges represent, whether the graphs are small or large, sparse or dense, and whether they are homogeneous or heterogeneous. We also survey relevant literature to identify key papers that have used these datasets and report their benchmark results for comparison.

This report presents our methodology, findings, and conclusions. It is structured as follows: Section 2 provides background on GNNs and related work; Section 3 introduces the datasets; Section 4 outlines the pipeline implementation; Section 5 presents the experiments and results; and Section 6 concludes the report and suggests possible improvements.

2 Methods

2.1 Datasets

In the context of graph classification, several properties are used to describe and compare graph datasets. These characteristics help determine how challenging a dataset is and what modeling strategies may be appropriate.

Small vs. Large Graphs This refers to the number of nodes and edges per individual graph in the dataset.

- **Small graphs** usually have tens of nodes and edges (e.g., MUTAG), making them computationally light and less structurally complex.
- **Large graphs** can have hundreds or thousands of nodes and edges, requiring more memory and processing power (e.g., datasets like REDDIT-BINARY or COLLAB).

Sparse vs. Dense Graphs This characterizes the ratio of the actual number of edges to the maximum possible number of edges.

- A graph is **sparse** if it has relatively few edges compared to the number of nodes (common in molecular or protein datasets).
- A graph is **dense** if it has many edges and a high degree of node connectivity.

Formally, sparsity can be approximated by the edge density: $\text{Density} = \frac{2 \cdot |E|}{|V| \cdot (|V| - 1)}$ where $|E|$ is the number of edges and $|V|$ is the number of nodes. With Density $D \geq 0.4$, a graph is denoted as dense in the following, a graph with $D \leq 0.1$ as very sparse and a graph with $D \approx 1$ as fully connected..

Homogeneous vs. Heterogeneous Graphs

- **Homogeneous graphs** contain only one type of node and one type of edge. Most classical GNN benchmarks (like MUTAG) fall in this category.
- **Heterogeneous graphs** include multiple types of nodes and/or edges, each potentially carrying different semantic information. For example, in biological datasets, nodes may represent different molecule components or protein substructures with diverse attributes.

Attributed vs. Non-Attributed Graphs

- In **attributed graphs**, nodes (and sometimes edges) carry feature vectors or labels, which can be used as input to GNN models.
- **Non-attributed graphs** only provide connectivity information, and node/edge features must be inferred or embedded manually.

Labeled vs. Unlabeled Graphs

- **Labeled graphs** provide a target class (e.g., "mutagenic" or "non-mutagenic") for each graph, necessary for supervised learning.
- **Unlabeled graphs** are more suitable for unsupervised or self-supervised learning approaches.

Multi-class vs. Binary Classification This refers to the number of target classes:

- **Binary classification** tasks (e.g., MUTAG, PROTEINS) distinguish between two labels.
- **Multi-class classification** tasks (e.g., ENZYMES) require assigning one label from more than two categories.

To evaluate the performance of Graph Neural Networks (GNNs) in graph classification tasks, we employ three well-established datasets from the TUDataset collection: **MUTAG**, **PROTEINS**, and **ENZYMES**. These datasets are drawn from domains such as chemistry and bioinformatics and are widely used in benchmarking GNN architectures due to their distinct structural properties and classification challenges. In the following sections, we present a detailed description of each dataset, including the semantics of nodes and edges, the size and density of the graphs, and the nature of their labels.

MUTAG

The MUTAG dataset consists of 188 chemical compounds, each modeled as an undirected graph. The classification task is to predict whether each compound exhibits mutagenic activity on a specific bacterium strain (*Salmonella typhimurium*). This is a binary classification problem.

Graph Structure: Each graph represents a molecule where:

- **Nodes** correspond to atoms in the molecule, such as Carbon (C), Oxygen (O), Nitrogen (N), etc. Each node is associated with a discrete label that identifies the atom type.
- **Edges** represent chemical bonds between atoms. These are undirected and unweighted, with no bond-type information explicitly provided in the basic version of the dataset.

Characteristics:

- **Number of graphs:** 188
- **Average number of nodes per graph:** 17.93
- **Average number of edges per graph:** 19.79
- **Mean Density:** $D \approx 0.096$
- **Classes:** 2 (mutagenic, non-mutagenic)
- **Type:** Small, very sparse, homogeneous, labeled, attributed

Scientific Usage: MUTAG has been used extensively as a benchmark in kernel-based and neural network-based graph learning models. Notable references include:

- Weisfeiler-Lehman subtree kernel [9]
- Graph Convolutional Networks (GCN) [5]
- Graph Isomorphism Network (GIN) [11]

PROTEINS

The PROTEINS dataset is derived from the protein data bank and contains 1,113 protein graphs. The classification task is to determine whether each protein is an enzyme or a non-enzyme, resulting in a binary classification problem.

Graph Structure: Each graph models the secondary structure of a protein:

- **Nodes** correspond to secondary structure elements (SSEs), such as alpha helices or beta sheets. Nodes are labeled with discrete types and may carry additional structural features.
- **Edges** connect SSEs that are in close spatial proximity or connected sequentially. These edges are undirected and encode structural relationships.

Characteristics:

- **Number of graphs:** 1,113
- **Average number of nodes per graph:** 39.06
- **Average number of edges per graph:** 72.82
- **Mean Density:** $D \approx 0.098$
- **Classes:** 2 (enzyme, non-enzyme)
- **Type:** Medium-sized, very sparse, heterogeneous, labeled, attributed

Scientific Usage: PROTEINS is a common benchmark in both classical and deep graph learning literature. It was used in:

- Deep Graph Kernels [12]
- Message Passing Neural Networks (MPNNs) [8]
- Geometric deep learning benchmarks [1]

ENZYMES

ENZYMES contains 600 protein graphs that are classified into one of six enzyme classes as defined by the Enzyme Commission (EC) number classification. This is a multi-class classification problem.

Graph Structure: Each graph represents a protein, modeled similarly to the PROTEINS dataset:

- **Nodes** denote SSEs in the protein structure, labeled by type and possibly enriched with additional feature information.
- **Edges** indicate spatial or sequential proximity between SSEs. These are undirected and generally sparse.

Characteristics:

- **Number of graphs:** 600
- **Classes:** 6 (corresponding to EC enzyme classes)
- **Average number of nodes per graph:** 32.63
- **Average number of edges per graph:** 62.14
- **Mean Density:** $D \approx 0.12$
- **Type:** Medium-sized, sparse, heterogeneous, labeled, attributed

Scientific Usage: ENZYMES presents a more complex classification challenge due to its multi-class nature and greater intra-class variability. It has appeared in:

- GIN benchmark experiments [11]
- TUDataset benchmarking [8]
- Comparative studies between graph kernels and deep models [12]

2.2 Summary Table

Dataset	# Graphs	Avg. Nodes	Avg. Edges	Classes	Graph Type ¹	Domain
MUTAG	188	17.93	19.79	2	S, SP, HO	Chemistry
PROTEINS	1113	39.06	72.82	2	M, SP, HE	Biology
ENZYMES	600	32.63	62.14	6	M, SP, HE	Biology

Table 1: Overview of the main properties of the selected TUDatasets.

(S=Small, SP=Sparse, HO=Homogenous, M=Medium, HE=Heterogenous)

2.3 Graph Neural Network Architectures

Graph Neural Networks (GNNs) are designed to process graph-structured data by leveraging the relationships (edges) between entities (nodes). The key idea is to iteratively update node representations by aggregating information from their neighbors. In this section, we discuss four prominent GNN architectures: Graph Convolutional Networks (GCN), Graph Isomorphism Networks (GIN), Graph Attention Networks (GAT), and GraphSAGE. Each of these models defines a different message passing and aggregation strategy.

2.4 Graph Convolutional Network (GCN)

Graph Convolutional Networks (GCNs) were introduced in [5] as an extension of the convolution operation to graph-structured data. GCNs perform a form of *neighborhood averaging*, where each node updates its representation based on the features of its neighbors and itself.

The forward propagation rule for a single GCN layer is:

$$\mathbf{H}^{(l+1)} = \sigma \left(\tilde{\mathbf{D}}^{-1/2} \tilde{\mathbf{A}} \tilde{\mathbf{D}}^{-1/2} \mathbf{H}^{(l)} \mathbf{W}^{(l)} \right) \quad (1)$$

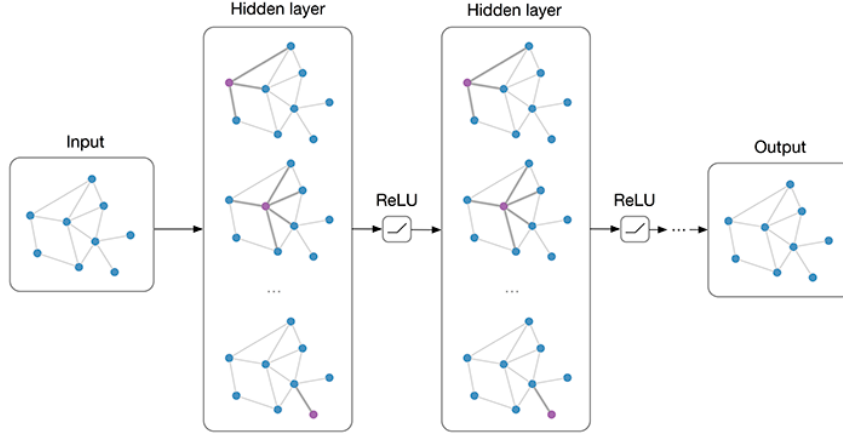


Figure 1: Graph Convolutional Neural Network [4].

Explanation of Symbols:

- $\mathbf{H}^{(l)}$ is the matrix of node features at layer l ; each row is the feature vector for one node.
- $\mathbf{H}^{(0)}$ is the input feature matrix (e.g., atom types).
- $\tilde{\mathbf{A}} = \mathbf{A} + \mathbf{I}$ is the adjacency matrix of the graph with added self-loops.
- $\tilde{\mathbf{D}}$ is the diagonal degree matrix of $\tilde{\mathbf{A}}$, where $\tilde{D}_{ii} = \sum_j \tilde{A}_{ij}$.
- $\mathbf{W}^{(l)}$ is a trainable weight matrix at layer l .
- σ is a non-linear activation function (typically ReLU).

The normalization $\tilde{\mathbf{D}}^{-1/2} \tilde{\mathbf{A}} \tilde{\mathbf{D}}^{-1/2}$ ensures that feature aggregation is scaled properly, especially for nodes with high degree. GCNs are efficient and perform well on many tasks, but they are limited in expressivity due to their linear nature and tendency to oversmooth node features across layers.

2.5 Graph Isomorphism Network (GIN)

The Graph Isomorphism Network (GIN), proposed in [11], is a GNN architecture specifically designed to match the discriminative power of the Weisfeiler-Lehman (WL) test – a classical test for graph isomorphism.

The update rule in GIN is:

$$\mathbf{h}_v^{(l+1)} = \text{MLP}^{(l)} \left((1 + \epsilon) \cdot \mathbf{h}_v^{(l)} + \sum_{u \in \mathcal{N}(v)} \mathbf{h}_u^{(l)} \right) \quad (2)$$

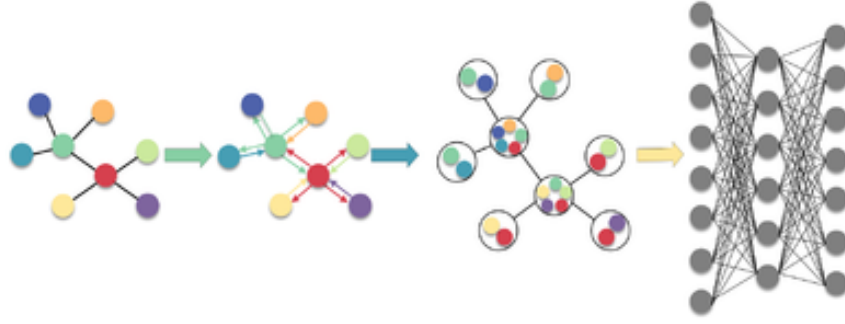


Figure 2: Node update process of a GIN layer [7].

Explanation of Symbols:

- $\mathbf{h}_v^{(l)}$ is the feature vector of node v at layer l .
- $\mathcal{N}(v)$ denotes the set of neighboring nodes of v .
- ϵ is a learnable or fixed scalar that determines the importance of the central node's own features.
- $\text{MLP}^{(l)}$ is a multi-layer perceptron applied after aggregation.

Unlike GCNs, which average the features, GIN performs a *sum aggregation* followed by a non-linear transformation, which retains more unique structural information. This design choice makes GIN theoretically more powerful for distinguishing non-isomorphic graphs. GIN performs particularly well in graph classification tasks such as MUTAG and ENZYMES.

2.6 Graph Attention Network (GAT)

Graph Attention Networks (GATs), introduced in [10], apply an attention mechanism to graphs. Instead of treating all neighbors equally, GAT learns to assign different importance (attention scores) to each neighbor.

The update rule for GAT is:

$$\mathbf{h}_v^{(l+1)} = \sigma \left(\sum_{u \in \mathcal{N}(v)} \alpha_{vu}^{(l)} \cdot \mathbf{W}^{(l)} \mathbf{h}_u^{(l)} \right) \quad (3)$$

Explanation of Symbols:

- $\mathbf{h}_v^{(l)}$ and $\mathbf{h}_u^{(l)}$ are the feature vectors of node v and its neighbor u .
- $\mathbf{W}^{(l)}$ is a trainable weight matrix.
- $\alpha_{vu}^{(l)}$ is the attention coefficient from node v to node u , learned via:

$$\alpha_{vu}^{(l)} = \frac{\exp \left(\text{LeakyReLU} \left(\mathbf{a}^T [\mathbf{W}^{(l)} \mathbf{h}_v^{(l)} \parallel \mathbf{W}^{(l)} \mathbf{h}_u^{(l)}] \right) \right)}{\sum_{k \in \mathcal{N}(v)} \exp \left(\text{LeakyReLU} \left(\mathbf{a}^T [\mathbf{W}^{(l)} \mathbf{h}_v^{(l)} \parallel \mathbf{W}^{(l)} \mathbf{h}_k^{(l)}] \right) \right)}$$

- \mathbf{a} is a trainable attention vector; \parallel denotes vector concatenation.




Figure 3: Graph Attention Mechanism [2].

By learning attention scores, GAT can focus on the most relevant parts of a node’s neighborhood, which is useful for datasets where not all connections are equally informative. It is particularly effective in social network data, though its performance in small molecular graphs is also competitive.

2.7 GraphSAGE

GraphSAGE (Graph Sample and Aggregate), introduced in [3], was designed for *inductive learning*, i.e., making predictions on unseen graphs. Unlike full-neighborhood aggregation, GraphSAGE samples a fixed-size neighborhood and aggregates features using various strategies (mean, max-pool, LSTM).

The generic update rule is:

$$\mathbf{h}_v^{(l+1)} = \sigma \left(\mathbf{W}^{(l)} \cdot \text{AGG}^{(l)} \left(\left\{ \mathbf{h}_v^{(l)} \right\} \cup \left\{ \mathbf{h}_u^{(l)} \mid u \in \mathcal{N}(v) \right\} \right) \right) \quad (4)$$

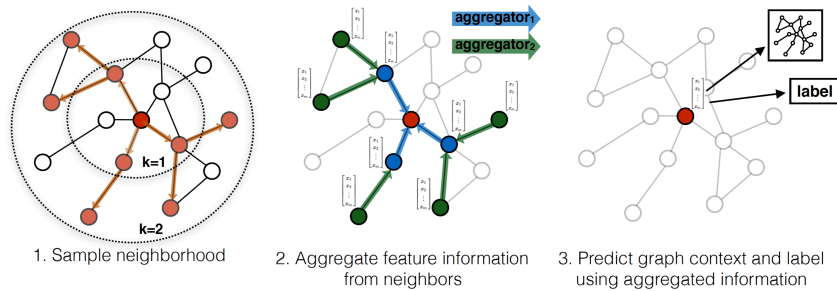


Figure 4: Graphsage Node update [6].

Explanation of Symbols:

- $\text{AGG}^{(l)}$ is an aggregation function (e.g., mean, LSTM, max-pooling).
- $\mathbf{h}_v^{(l)}$ is the current node embedding at layer l .
- $\mathbf{W}^{(l)}$ is a trainable weight matrix.
- σ is a non-linear activation function.

GraphSAGE is particularly scalable for large graphs, such as citation or social networks, due to its sampling approach. It also enables efficient inference on previously unseen nodes and graphs

3 Training Setup

3.1 Data Preprocessing

For all experiments, we used datasets from the TUDataset collection, loaded via a custom `TUDatasetLoader` interface. Each dataset was preprocessed using a consistent set of transformations:

- **ToUndirected:** All graphs were converted to undirected format, ensuring edge symmetry.
- **NormalizeFeatures:** Node features were scaled to have unit norm to stabilize training.
- **RemoveIsolatedNodes:** Any node with no neighbors was removed to avoid redundant computation and disconnected components.

Node attributes were explicitly used via `use_node_attr=True`, which is critical for datasets like ENZYMES and PROTEINS that rely on rich node-level descriptors.

3.2 Model Architecture

All GNN variants followed a modular architecture comprising two main components:

- **Backbone:** A stack of three GNN layers specific to the chosen model type—GCN, GAT, GIN, or GraphSAGE. All of the used models were custom engineered versions of the original models, offering high modularity and controllability by relying on completely configurable blocks.
- **Classification Head:** A fully connected MLP of varying depth, applied after an aggregating pooling operation, mapping the graph embedding to logits for classification.

This design ensured consistency across models, allowing a fair comparison of the expressive capacity of the GNN layers themselves.

3.3 Optimization and Training Strategy

We used the Adam optimizer with specific learning rates and weight decays per dataset. The values were achieved by trying different configuration alongside different learning rates, thus followed an handcrafted hyperparameter-optimization. The learning rate was decayed exponentially by a factor of $\gamma = 0.95$ at each epoch using the `ExponentialLR` scheduler. These hyperparameters were chosen based on stability across datasets and prior work in GNN literature.

3.4 Cross-Validation and Early Stopping

To ensure robustness of results and handle limited dataset sizes, we adopted a 10-fold stratified cross-validation approach using a custom `KFoldRunner`. The key parameters were:

- **Number of Folds:** 10 (with validation performed on each fold sequentially).
- **Batch Size:** 16 for training; 1 for validation and testing.
- **Patience:** Training was monitored via validation accuracy, with early stopping triggered if no improvement larger than `abort condition` was seen for 50 epochs.

3.5 Evaluation Metric

Validation accuracy (`val_acc`) was used as the primary performance metric, with the goal of maximizing it. Mean and standard deviation across folds were computed and reported for each model-dataset combination. All the tested configs are available in the GitHub repo.

4 Results and Discussion

We performed a comparative evaluation of four prominent Graph Neural Network (GNN) architectures—Graph Attention Network (GAT), Graph Convolutional Network (GCN), Graph Isomorphism Network (GIN), and GraphSAGE—on three graph classification datasets from the TUDataset benchmark collection: MUTAG, PROTEINS and ENZYMES.

For each model-dataset combination, we report validation accuracy obtained over multiple training runs using a consistent pipeline. The performance distributions are visualized using boxplots.

4.1 Performance of GNNs Across Datasets

GAT Performance

As shown in Figure 5, GAT performs well on the MUTAG dataset, achieving an average accuracy above 0.8, with low variance. This suggests that attention mechanisms effectively capture relevant local structures in small molecular graphs. However, performance drops significantly

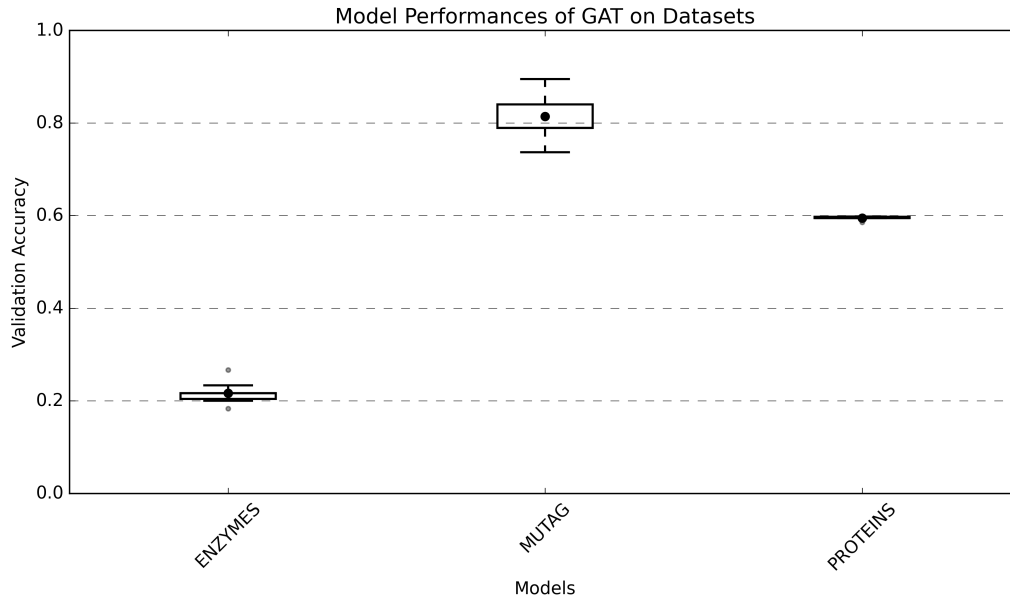


Figure 5: Validation accuracy of GAT across ENZYMES, MUTAG, and PROTEINS datasets.

on ENZYMES (accuracy around 0.2), indicating that GAT struggles with more complex, multi-class, heterogeneous biological graphs. On PROTEINS, performance remains modest (around 0.6), suggesting limited effectiveness of attention-based aggregation in this domain.

GCN Performance

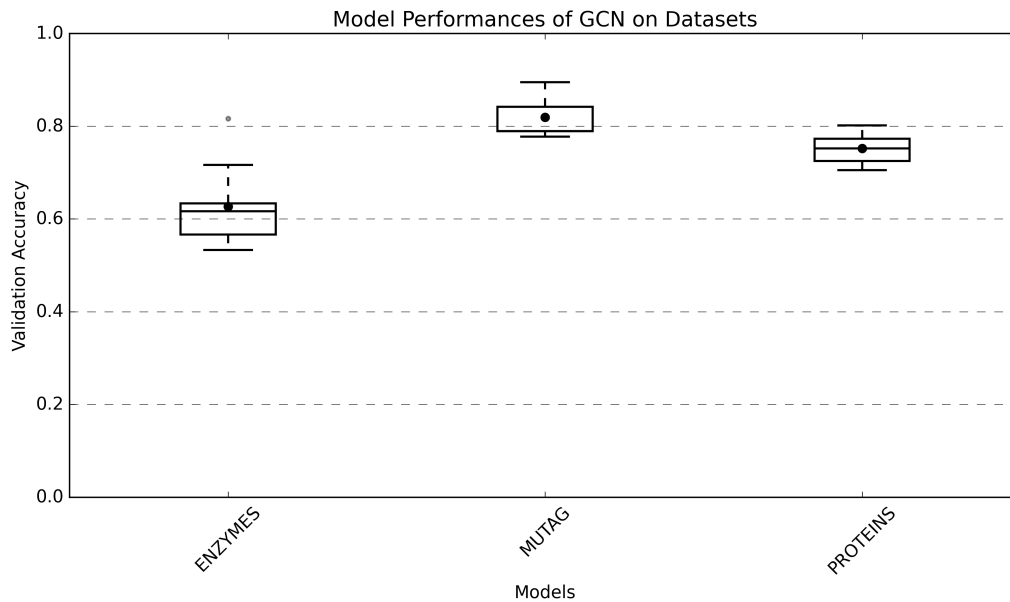


Figure 6: Validation accuracy of GCN across ENZYMES, MUTAG, and PROTEINS datasets.

GCN (Figure 6) consistently performs well across all datasets. On ENZYMES, it outperforms all other models with a mean accuracy around 0.65, demonstrating strong generalization on complex graphs. It also maintains robust accuracy on MUTAG and PROTEINS, both above 0.75, confirming its adaptability and simplicity as a strong baseline.

GIN Performance




Figure 7: Validation accuracy of GIN across ENZYMES, MUTAG, and PROTEINS datasets.

GIN (Figure 7) performs exceptionally well on MUTAG (up to 0.9), leveraging its theoretical expressiveness. It also performs competitively on PROTEINS, but yields slightly lower and more variable results on ENZYMES. This suggests that GIN may benefit from richer node labels and simpler structural motifs, which are less prominent in ENZYMES.

GraphSAGE Performance




Figure 8: Validation accuracy of GraphSAGE across ENZYMES, MUTAG, and PROTEINS datasets.

GraphSAGE (Figure 8) demonstrates strong and consistent performance across all datasets. It

particularly excels on ENZYMES and PROTEINS, where it outperforms GIN and GAT. Its inductive sampling mechanism appears advantageous for generalizing over complex and variable topologies.

4.2 Model Comparison on Individual Datasets

ENZYMES Dataset

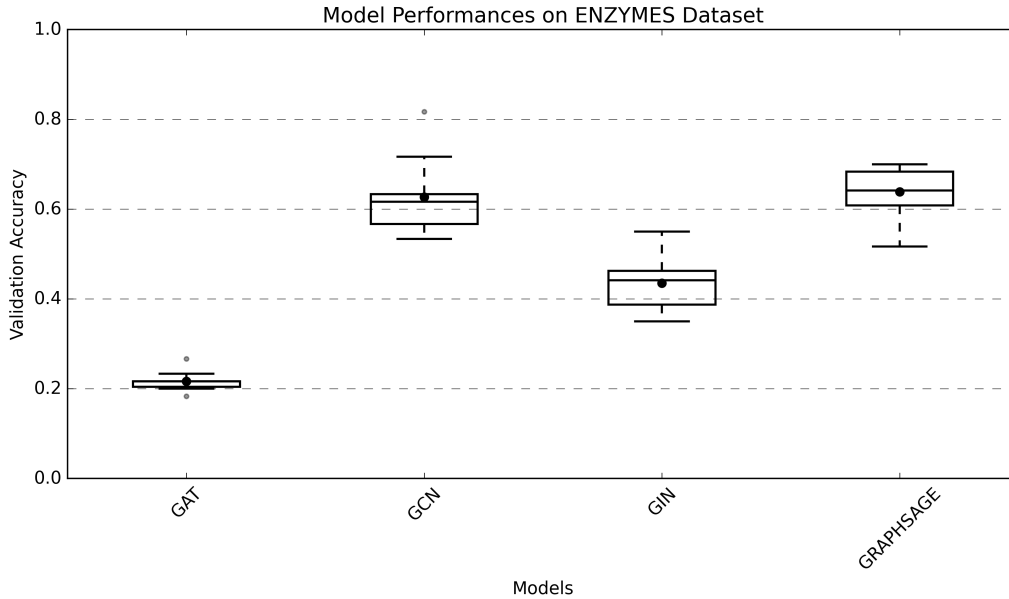


Figure 9: Comparison of GNN models on the ENZYMES dataset.

The ENZYMES dataset is multi-class and structurally heterogeneous, making it the most challenging. As shown in Figure 9, GCN and GraphSAGE outperform the others, reaching accuracy values above 0.6. GAT performs poorly (0.2), suggesting that it may not adapt well to biological networks with irregular connectivity and weak local regularities. GIN provides mid-level accuracy but with higher variance, indicating sensitivity to initialization or optimization.

MUTAG Dataset

MUTAG is a small binary classification dataset composed of molecular graphs. All models achieve strong results (Figure 10), with GIN and GCN slightly ahead, suggesting that most GNN architectures can effectively learn over this simpler domain. GraphSAGE, although consistent, slightly underperforms compared to the rest, possibly due to limited neighborhood size in small graphs.

PROTEINS Dataset

In Figure 11, we observe that GCN, GIN, and GraphSAGE all achieve comparable performance (0.750.8), while GAT again falls behind. The consistent performance of GCN and GraphSAGE confirms their generalization ability, while GINs strong results again demonstrate its suitability

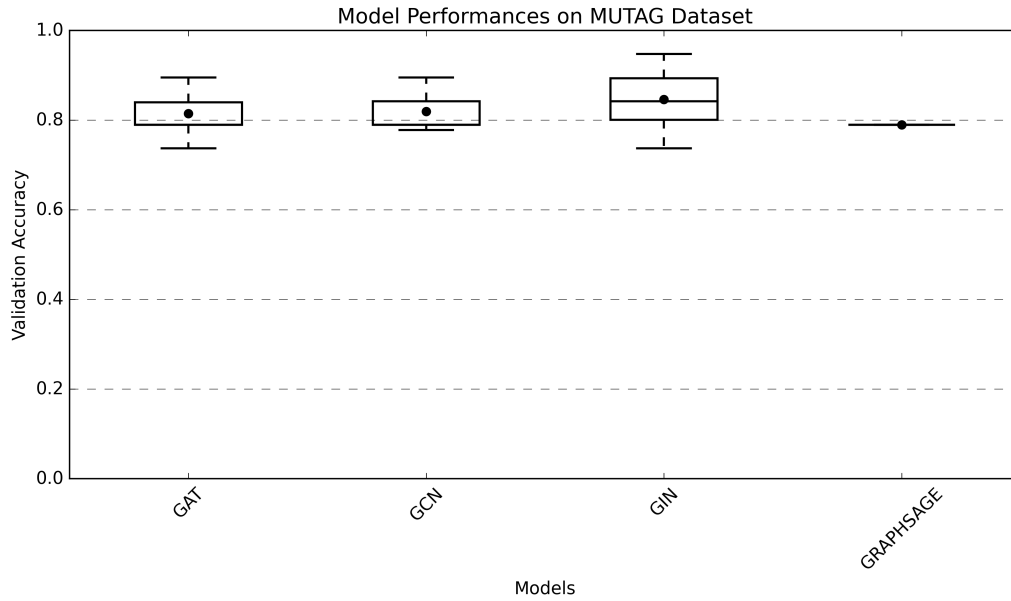


Figure 10: Comparison of GNN models on the MUTAG dataset.

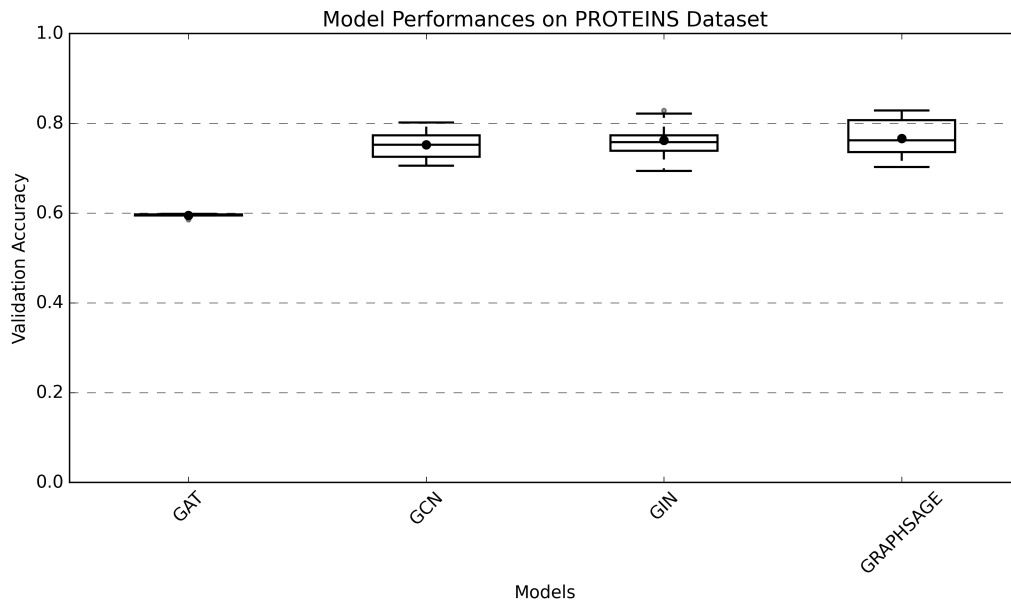


Figure 11: Comparison of GNN models on the PROTEINS dataset.

for structural pattern learning. GAT’s weakness suggests its attention mechanism may not align well with the topological characteristics of protein interaction graphs.

4.3 General Observations

- **GCN** shows excellent performance and generalization across all datasets. Its simplicity and effectiveness make it a strong baseline.
- **GraphSAGE** performs robustly, particularly in large, noisy, or irregular graphs (ENZYMES, PROTEINS), benefiting from its sampling-based inductive design.
- **GIN** shines in small, structurally simple datasets like MUTAG, but is more sensitive in

complex settings.

- **GAT** suffers in generalization and is highly dataset-dependent. While it performs well on MUTAG, its attention mechanism appears to overfit or fail in more complex domains.

These results suggest that no single GNN architecture dominates in all settings. Instead, model performance is closely tied to the nature of the graph data—node labels, class structure, graph size, and connectivity patterns. Hence, model selection should consider dataset characteristics.

References

- [1] Vijay Prakash Dwivedi, Chaitanya K Joshi, Thomas Laurent, Yoshua Bengio, and Xavier Bresson. Benchmarking graph neural networks. *arXiv preprint arXiv:2003.00982*, 2020.
- [2] Giuseppe Fùtia. Graph attention networks under the hood, 2021.
- [3] William L Hamilton, Rex Ying, and Jure Leskovec. Inductive representation learning on large graphs. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2017.
- [4] Thomas Kipf. Graph convolutional networks, 2016.
- [5] Thomas N Kipf and Max Welling. Semi-supervised classification with graph convolutional networks. *International Conference on Learning Representations (ICLR)*, 2017.
- [6] Jure Leskovec. Graphsage: Inductive representation learning on large graphs, 2017.
- [7] Tong Lu, Sizu Hou, and Yan Xu. Fault line selection algorithm for distribution networks based on adapglin network. *IET Generation, Transmission Distribution*, 2023.
- [8] Christopher Morris, Martin Ritzert, Matthias Fey, William L Hamilton, Jan Eric Lenssen, Gaurav Rattan, and Martin Grohe. Tudataset: A collection of benchmark datasets for learning with graphs. *arXiv preprint arXiv:2007.08663*, 2020.
- [9] Nino Shervashidze, Pascal Schweitzer, Erik Jan van Lee, Kurt Mehlhorn, and Karsten M Borgwardt. Weisfeiler-lehman graph kernels. In *Advances in Neural Information Processing Systems*, volume 24, 2011.
- [10] Petar Velickovi, Guillem Cucurull, Arantxa Casanova, Adriana Romero, Pietro Lio, and Yoshua Bengio. Graph attention networks. In *International Conference on Learning Representations (ICLR)*, 2018.
- [11] Keyulu Xu, Weihua Hu, Jure Leskovec, and Stefanie Jegelka. How powerful are graph neural networks? In *International Conference on Learning Representations (ICLR)*, 2019.
- [12] Pinar Yanardag and S V N Vishwanathan. Deep graph kernels. In *Proceedings of the 21th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pages 1365–1374, 2015.